



Time to First Token



Time per Output Token



End-to-End Latency

